




LOCAL-AGENT READINESS • RYZEN AI MAX+ 395

Almost There.

Local coding agents on an AMD AI PC — a measured readiness report

Harsh Singh — Research & Advanced Development, AMD

Supervisors: Eddie Richter • Paul Hartke



Contents

Executive summary

1 • Why local coding agents, and why AMD

2 • Background & related work

3 • The platform

4 • Methodology & data hygiene

5 • The model lineup

6 • Headline capability

7 • Per-run results

8 • Capability vs the frontier

9 • Why the models perform as they do

10 • Failure-mode taxonomy

11 • Behavioral patterns

12 • Token efficiency

13 • Task difficulty

14 • Efficiency & intelligence per watt

15 • Model-specific profiles

16 • Infrastructure & integration failures

17 • Key findings & recommendations

18 • Limitations

References

Executive summary

We built a reproducible evaluation platform and ran local coding agents at scale on a single AMD Ryzen AI MAX+ 395: **848 trials, 10,355 agent steps, and 240.8 million tokens processed (3.84M generated)** across **14 models** and **4 agent harnesses**. After filtering infrastructure and agent-integration failures, **561 trials** form the analysis set behind every number in this report.

The central finding is hopeful but honest. The best on-device configuration — **terminus-2 driving Qwen3.5-35B-A3B** — solves **84% of the easy agentic-terminal slice** (21/25), a 35-billion-parameter *sparse* model that activates only ~3B parameters per token and runs entirely on a laptop-class APU. That is genuinely useful work. But it is the easy slice: on the full Terminal-Bench 2.0 leaderboard the best open-weight model (GLM-5, 52.4%) trails the frontier (Claude Opus 4.7, 90.2%) by ~38 points — yet under a *standardized* SWE-bench harness the open-vs-frontier gap is only ~1 point (76.8% vs 75.8%). The deficit is concentrated in long-horizon *agentic discipline* — loop control, output format, and knowing when you are actually done — not in raw coding ability.

848 trials run	561 valid	240.8M tokens processed	3.84M generated
14 models	4 harnesses	84% best local config	×5.3 intelligence/watt (2yr)

What the data says, and why.

- **The agent harness matters as much as the model.** Holding Qwen3.5-35B-A3B fixed, pass rate is **84% under terminus-2, 53% under openhands, 42% under qwen-coder** — a ~2× spread on identical weights. This is the agent-computer-interface thesis (Yang et al. 2024) reproduced on-device: the scaffold, not just the model, determines success.
- **Sparsity is the on-device unlock — and it is an architectural choice, not luck.** Sparse Mixture-of-Experts models (~3–5B active parameters) generate at **43–72 tok/s**; dense models of comparable quality crawl at **12–27 tok/s**. Decoupling *active* compute from *total* capacity (Shazeer et al. 2017; Fedus et al. 2022; Jiang et al. 2024) is exactly the trade AMD’s unified memory rewards.
- **Format adherence is trained, not emergent.** Coder/agent-tuned models (the Qwen3-Coder family) emit the structured action protocol cleanly; a dense reasoning model trained on textbooks (phi-4) wraps its JSON in prose and loses the turn. Malformed actions afflict **12%** of valid trials, concentrated in models never post-trained for tool use.
- **Failure is about discipline, not capability.** The dominant behavioral failure is the **doom-loop (21% of valid trials)** — repetition until truncation, the agentic echo of neural-text degeneration (Holtzman et al. 2020) — and failing runs *flail* rather than recover, consistent with the limits of intrinsic self-correction (Huang et al. 2024): they spend **3–8× more tokens** than passing runs without converging.
- **“GLM-4.7” is two different animals.** The cloud model (~32B active) scores 77%; the local GLM-4.7-Flash (~3B active, Q4-quantized) scores 47% — a ~10× active-capacity gap compounded by low-bit quantization that disproportionately harms instruction-following.
- **Local efficiency is competitive, and improving fast.** Generation efficiency lands at **~0.66–1.1 tokens/joule** on the AMD power envelope; intelligence-per-watt rose **5.3×** in two years (Saad-Falcon et al. 2025). Our snapshot is one point on a steep curve.

1 • Why local coding agents, and why AMD

Inference is migrating from the data center to the device. The *Intelligence per Watt* study (Saad-Falcon et al. 2025, Stanford Hazy Research + Together AI, arXiv:2511.07885) frames this as a shift from a “mainframe era” to a “PC era” of AI and quantifies it: local models now accurately answer **88.7%** of single-turn queries; local query coverage climbed from **23.2% → 71.3%** between 2023 and 2025; and intelligence per watt — accuracy per unit of power — improved **5.3×** over the same window (3.1× from models, 1.7× from hardware). A Qwen3-32B on consumer silicon runs at only **1.5× lower IPW than an NVIDIA B200** on the same model.

For AMD this is strategic. The Ryzen AI MAX+ 395 (“Strix Halo”) pairs 16 Zen 5 cores, a 40-CU RDNA 3.5 iGPU, and a 50-TOPS XDNA2 NPU with **up to 128 GB of unified LPDDR5X** — enough to hold models a discrete consumer GPU cannot, and, critically, to keep an entire Mixture-of-Experts parameter pool resident in shared memory so that expert routing incurs no PCIe transfer. AMD’s own framing is that the PC is becoming “an intelligent assistant that works alongside you,” with on-device inference delivering “data security, privacy, increased responsiveness, and the ability to use applications even when not connected to the internet.” A coding *agent* — which reads a repository, runs commands, and iterates toward a goal — is the most demanding and most valuable form of that assistant. This report asks the concrete question: **can today’s local models be that agent on AMD hardware, and where exactly do they fall short?**

2 • Background & related work

This work sits at the intersection of six literatures; each supplies a lens we use to interpret our measurements.

Agentic coding benchmarks have bifurcated, and that split is our measurement frame. SWE-bench (Jimenez et al. 2024) asks whether a model can *write the patch* for a real GitHub issue, and its headline that the best 2023 model resolved under 2% of tasks established the modern “agents on real software” framing; SWE-bench Verified (OpenAI 2024) is the human-validated 500-task subset that became the frontier coding scoreboard. Terminal-Bench (Merrill et al. 2026) — the benchmark we use — measures something orthogonal: whether an agent can *drive a long-horizon loop* in a real shell, and it reports frontier agents below 65% while explicitly recording *which harness* produced each score. This “write the fix vs. drive the loop” split is why our 84% on an easy slice and the large remaining frontier gap are both legible: Terminal-Bench isolates loop-driving competence from patch-writing.

The scaffold rivals the model. A consistent thread runs from ReAct (Yao et al. 2023), which established the interleaved reason→act→observe loop nearly every coding agent now uses, through Reflexion (Shinn et al. 2023), which added weight-free verbal self-reflection to recover from failed trials, to SWE-agent (Yang et al. 2024), whose agent-computer-interface (ACI) thesis showed that — *holding the model fixed* — a purpose-built interface drives large gains. OpenHands (Wang et al. 2024), one of the harnesses we benchmark, generalizes this into an open platform. Our same-model 84/53/42% spread across terminus-2, OpenHands, and qwen-coder is fresh, on-device evidence for the ACI thesis.

Failure mechanisms are known. Holtzman et al. (2020) showed that likelihood-maximizing decoding drives models into degenerate repetition — the mechanistic basis for our dominant doom-loop. Huang et al. (2024) demonstrated that LLMs cannot reliably self-correct reasoning without external feedback and sometimes degrade when they try — which explains why our failing agents flail (3–8× the tokens, far more steps) rather than escape a bad trajectory. Together these reframe the bottleneck as *agentic discipline*, not raw coding ability.

Small and open models have become genuinely capable, the premise of running a 35B-class model on an APU at all. Phi-3 (Abdin et al. 2024) showed a 3.8B model rivaling Mixtral-8×7B quality on a phone; Qwen2.5-Coder (Hui et al. 2024) reported state-of-the-art-for-size coding; Llama 3 (Grattafiori et al. 2024), gpt-oss (OpenAI 2025), and GLM-4.5 (Zeng et al. 2025) brought open agentic and explicitly MoE designs into the open-weights tier in exactly the sparse active-parameter regime our efficiency results exploit.

Mixture-of-Experts is the efficiency engine. Shazeer et al. (2017) introduced the sparsely-gated MoE layer that decouples total capacity from per-token compute; Switch Transformer (Fedus et al. 2022) formalized “active ≠ total” by scaling parameters at near-constant per-token FLOPs; Mixtral (Jiang et al. 2024) made it concrete and open (47B total / ~13B active). This is precisely why our sparse models (~3B active) run 3–5× faster than dense peers, and why unified memory — which holds the whole expert pool with no transfer cost — is the right substrate.

On-device deployment and quantization complete the picture. Edge-LLM surveys (Zheng et al. 2024) catalogue the compression-and-co-design pipeline that makes local agents feasible; controlled quantization studies show FP8 is effectively lossless (Kurtic et al. 2024, “Give Me BF16 or Give Me Death?”) while low-bit integer quantization (the Q4_K_M GGUF format most of our runs used) costs a few points on average but disproportionately harms instruction-following and multi-step reasoning — exactly what an agent

benchmark stresses. The *Intelligence per Watt* framing (Saad-Falcon et al. 2025) is the efficiency metric we extend from single-turn QA to multi-step agentic tasks on a specific AMD accelerator.

Our contribution is the combination in one place: a filtered, on-device, AMD-specific evaluation that pairs a concrete capability result with a same-model cross-harness ablation *and* an architecture-grounded behavioral failure taxonomy showing the local-agent bottleneck is discipline, not coding ability.

3 • The platform

The first deliverable is a platform you can run. A single config-driven runner (`bin/sweep.sh`) drives the `harbor` CLI (Terminal-Bench 2.0) across a list of models and tasks; a model-lifecycle engine (`lib/model_manager.sh`) downloads each GGUF model, serves it via `llama-server`, warms it, and tears it down between models. Agent harnesses are declarative profiles (`config/agents/*.conf`) encoding each agent's networking mode and flags, so a sweep is one command and a new agent is one file:

```
./bin/sweep.sh -a terminus-2 -m config/models/batch1.txt -t config/tasks/easy.txt
```

A trial passes when the task's verifier writes `reward.txt = 1`. Crucially, **every LLM call is logged with full token usage and llama.cpp timings** (generation tok/s, prefill tok/s, cache hits, finish reason) — which is what makes the validity filtering, token-efficiency, and per-watt analysis in this report possible. An interactive dashboard (`dashboard/index.html`) lets a reader explore capability, efficiency, failure modes, and individual trajectories down to single generations.

Scope of the study. Across the campaign we recorded **848 trials**, **10,355 agent steps**, and **240.8M tokens processed (3.84M generated)** spanning **14 models** and **4 agent harnesses** — all on one AMD APU.

4 • Methodology & data hygiene

Raw pass rates conflate model behavior with infrastructure noise. We classify every trial and **use only valid trials for every number in this report**; dropped trials are fully accounted for (§16) but never reported as capability findings.

REASON	TRIALS	RULE
valid	561	used for all capability / efficiency / failure analysis
agent_integration_failure	172	swe-agent / mini-swe-agent: harness crashed before real work
ungraded	101	no verifier reward written (environment/agent error mid-run)
infra_slow	8	generation tok/s < max(5, 0.4 × the model's own baseline) = server stall
test_run	6	test / probe runs

(848 = 561 + 172 + 101 + 8 + 6. A further 91 trial rows from an early `/var/tmp/harbor-results` jobs-dir were **deduplicated** before this accounting — verified to be byte-identical copies of trials already present under the named runs, so counting them would have double-counted three models.)

The key discriminator is **throughput relative to each model's own baseline**. A model that is simply slow (a dense 22B at ~12 tok/s) keeps its low-but-consistent baseline and stays valid; a run whose throughput *collapses* to a fraction of that model's norm (e.g. a GLM-4.7-Flash trial at **1.2 tok/s** against a **45.8** baseline) is a server stall and is dropped. Two harnesses (`swe-agent`, `mini-swe-agent`) are excluded wholesale: their trajectories show the harness crashing before the model does any real work (a literal unexpanded `$(pwd)` → `NoSuchPathError`; runs that never reach an assistant turn). That is harness immaturity, not model capability — a distinction that matters for where engineering effort should go.

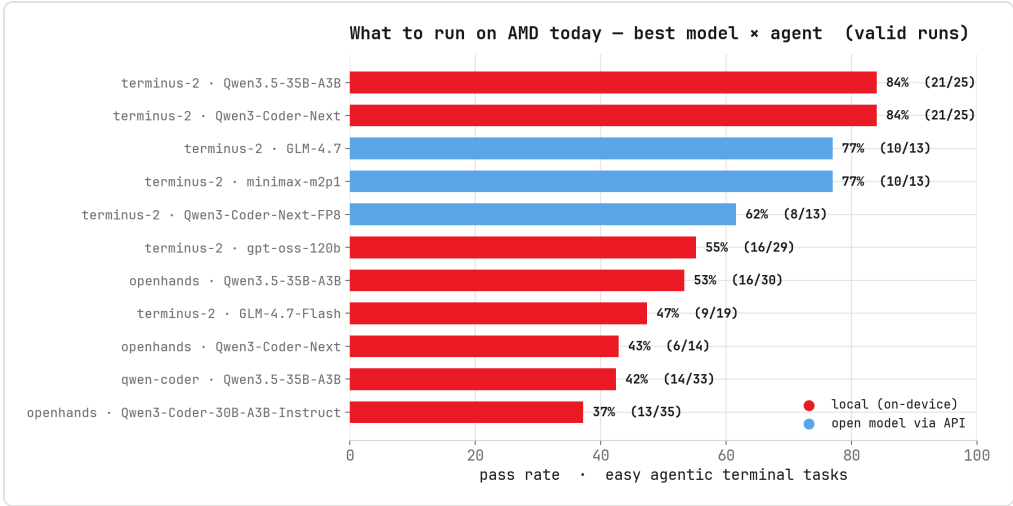
5 • The model lineup

Fourteen models span four labs’ design philosophies and two architecture families. Because the rest of this report explains behavior in terms of these choices, we tabulate them first. “Active” is the parameters used per token (the inference-cost driver); “total” is the capacity.

MODEL	LAB	ARCH	TOTAL	ACTIVE	TRAINING FOCUS
Qwen3.5-35B-A3B	Alibaba (Qwen)	sparse MoE	~35B	~3B	general
Qwen3-Coder-Next	Alibaba (Qwen)	sparse MoE (non-reasoning)	~80B	~3B	coder / agentic
Qwen3-Coder-30B-A3B-Instruct	Alibaba (Qwen)	sparse MoE	~30B	~3B	coder / agentic
GLM-4.7 (cloud)	Zhipu / Z.ai	sparse MoE	~355B	~32B	reasoning / coding / agentic
GLM-4.7-Flash (local)	Zhipu / Z.ai	sparse MoE	~30B	~3B	coder / agentic (local)
MiniMax-M2.1	MiniMax	sparse MoE	~230B	~10B	coder / agentic
gpt-oss-120b	OpenAI	sparse MoE (MXFP4)	117B	5.1B	reasoning / agentic
gpt-oss-20b	OpenAI	sparse MoE (MXFP4)	21B	3.6B	reasoning / agentic
Nemotron-3-Nano-30B-A3B	NVIDIA	hybrid Mamba-2 + MoE	~32B	~3B	reasoning (emits traces)
phi-4	Microsoft	dense	14B	14B	reasoning (synthetic/textbook)
Mistral-Small-2409	Mistral AI	dense	22B	22B	general + function-calling
Llama-3.1-8B-Instruct	Meta	dense	8B	8B	general
gemma-3-12b-it	Google	dense	~12B	~12B	general / multimodal
Qwen3-Coder-Next-FP8	Alibaba (Qwen)	sparse MoE (FP8)	~80B	~3B	coder / agentic

Most local runs used **Q4_K_M GGUF** quantization (the memory-efficient default); the one FP8 variant is called out because its near-lossless quantization isolates a quantization effect from an architecture effect (§11).

6 • Headline capability



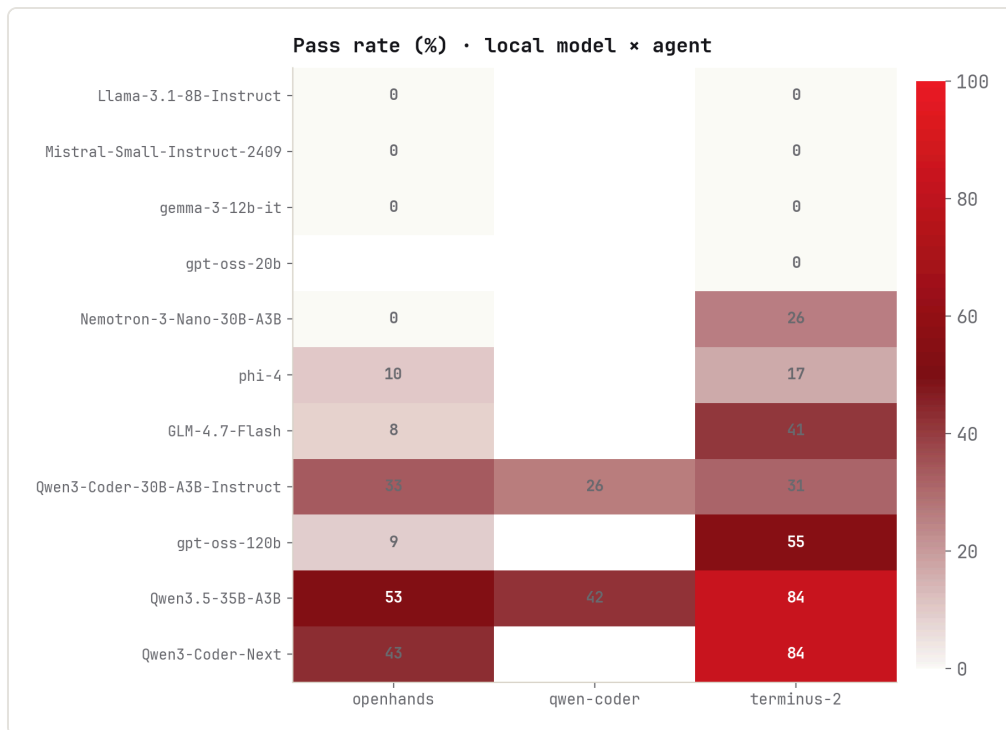
Best model × agent configurations by pass rate on the easy agentic-terminal slice (valid runs). Local configs in red; open models served via API in blue.

Two configurations top the board at **84%** (21/25): terminus-2 with **Qwen3.5-35B-A3B** and terminus-2 with **Qwen3-Coder-Next**. The aggregate picture by harness:

Agent comparison

AGENT	VALID TRIALS	PASSES	PASS RATE	AVG STEPS	AVG COMP TOK	AVG MODEL-TIME
terminus-2	255	120	47%	20.4	8,682	163 s
qwen-coder	67	23	34%	—	—	—
openhands	239	47	20%	21.6	6,817	153 s

terminus-2 more than doubles openhands’ pass rate on the same hardware, and the gap is *not* a model gap — it is the same models behind a different interface. Two mechanisms explain it. First, openhands absorbs more integration failures (§16): its richer tool surface (a bash executor, a file editor, a browser) is more sensitive to a model’s chat-template and tool-schema conformance, so weaker or mis-templated models stall at zero steps where terminus-2’s single bash channel still engages. Second, even when both engage, terminus-2’s narrow bash-only loop gives the model fewer ways to go wrong; openhands’ breadth invites the model to wander between tools. This is the agent-computer- interface result (Yang et al. 2024) reproduced locally: a constrained, well-matched interface beats a powerful but loosely-coupled one, and the effect is sharpest on a single model — Qwen3.5-35B-A3B scores **84% / 53% / 42%** under terminus-2 / openhands / qwen-coder.



Pass rate by local model × agent harness. The terminus-2 column dominates; weaker harnesses suppress otherwise-capable weights.

What an 84% run actually looks like — a real `fix-git` trajectory, lightly trimmed:

```

● fix-git · terminus-2 · Qwen3.5-35B-A3B REWARD 1.0 · PASS

ep 3    "The cherry-pick encountered a merge conflict in _includes/about.md."
$
cat _includes/about.md

ep 5    "It failed because no editor is available to write the commit message."
$
git cherry-pick -continue -m "Move to Stanford"

ep 6    "The -m flag is being interpreted as -mainline for merge commits. I need to abort and use a non-interactive editor."
$
git status ; GIT_EDITOR=true git cherry-pick -continue

ep 7    "The cherry-pick completed. Verify the log and working tree."
$
git log -oneline -5 ; git status

ep 8    "The task is complete. The 'Move to Stanford' commit (cd2d6e1) is now HEAD; the tree is clean."
is_task_complete: true

Converged in 9 turns · recovered from three distinct obstacles · reward 1.0

```

Agent trace — a convergent solve. Qwen3.5-35B-A3B works through a merge conflict, an editor failure, and the `-m/-mainline` gotcha in nine tight turns, verifying before it declares done. Contrast the loops in §10.

The reasoning on display here is genuine, not pattern-matched: the model recognizes that `-m` is being parsed as `--mainline`, and works around the missing editor with `GIT_EDITOR=true`. The same adaptive problem-solving recurs across the lineup — and gpt-oss-120b even writes a complete, compiling Coq proof (§9.4). The strengths and the failure modes are two sides of the same models; this report documents both.

```
● broken-python · terminus-2 · Qwen3.5-35B-A3B REWARD 1.0 · PASS

ep 1    "The issue is clear: Python 3.13 is installed but the pip module is missing. The pip3 script exists but
        cannot run."
        $
        python3 -m ensurepip --upgrade

ep 4    "curl is not available in this environment. I need an alternative way to download get-pip.py – I can use
        Python's own urllib."
        $
        python3 -c "import urllib.request;
        urllib.request.urlretrieve('https://bootstrap.pypa.io/get-pip.py', 'get-pip.py')"

ep 6    "The pip installation has been repaired – both pip3 --version and python3 -m pip --version now work."
        is_task_complete: true

Diagnosed → adapted around a missing tool → repaired in 7 turns · reward 1.0
```

Agent trace – adaptive problem-solving. Qwen3.5-35B-A3B repairs a broken `pip`: it diagnoses the missing module, and when `curl` is unavailable it pivots to fetching the installer with Python's own `urllib` – real tool knowledge, the upside the failure traces in §10 should be read against.

7 • Per-run results

Every valid run, by harness. “Top failure” is the most common terminal state among that run’s failed trials.

terminus-2

RUN	N	PASS	PASS%	AVG STEPS	TOP FAILURE
qwen3-coder-next-64gb	25	21	84%	16.8	ran without solving
qwen3.5-35b-a3b-local	25	21	84%	12.2	ran without solving
glm47-together	13	10	77%	31.3	ran without solving
m2.1-fireworks	13	10	77%	44.4	ran without solving
qwen-next-together	13	8	62%	71.2	ran without solving
gpt-oss-120b-64gb	29	16	55%	4.1	ran without solving
glm-4.7-flash-local	19	9	47%	11.6	ran without solving
qwen-30b-local †	11	4	36%	26.1	ran without solving
nemotron-3-nano-64gb	23	7	30%	10.2	ran without solving
glm47-flash-local-v2	10	3	30%	19.7	ran without solving
qwen3-coder-30b-local	31	9	29%	26.3	ran without solving
phi4-14b	6	1	17%	7.0	ran without solving
nemotron-3-nano-local	8	1	12%	15.4	ran without solving
gpt-oss-20b-local	8	0	0%	2.5	ran without solving
llama31-8b-instruct	8	0	0%	26.9	ran without solving
mistral-small-2409	12	0	0%	21.0	ran without solving
gemma3-12b-it	1	0	0%	39.0	ran without solving

† `qwen-30b-local` and `qwen3-coder-30b-local` are two separate sweep runs of the **same** model (Qwen3-Coder-30B-A3B-Instruct); both are shown to expose run-to-run variance (36% vs 29%). Likewise, `qwen-next-together` (and `qwen-next-together-v2`, openhands) is the **Qwen3-Coder-Next-FP8** model served via the Together API — the §6 chart labels by model, this table by run. A model appearing in more than one run is why the per-run table has more rows than the 14-model lineup. (These 17 rows sum to terminus-2’s 255 valid trials.)

openhands

RUN	N	PASS	PASS%	AVG STEPS	TOP FAILURE
qwen3.5-35b-a3b-local	30	16	53%	19.5	ran without solving
qwen3-coder-next-64gb	14	6	43%	23.0	ran without solving
qwen3-coder-30b-local	35	13	37%	25.5	ran without solving
glm47-together-v2	13	4	31%	62.6	ran without solving
qwen-30b-local	13	3	23%	41.0	ran without solving
phi4-14b	10	1	10%	16.7	ran without solving
gpt-oss-120b-64gb	34	3	9%	2.7	no agent steps
glm47-flash-local-v2	12	1	8%	25.0	ran without solving
gemma3-12b-it	12	0	0%	0.0	no agent steps
llama31-8b-instruct	12	0	0%	42.5	ran without solving
m2.1-fireworks-v2	13	0	0%	0.0	no agent steps
mistral-small-2409	12	0	0%	1.7	no agent steps
nemotron-3-nano-64gb	16	0	0%	6.6	ran without solving
qwen-next-together-v2	13	0	0%	62.3	ran without solving

The four “no agent steps” openhands runs are integration failures (§16) — note the same models pass under terminus-2 (gpt-oss-120b: 55% vs 9%), the cleanest possible demonstration that a 0% is often a harness verdict, not a model verdict.

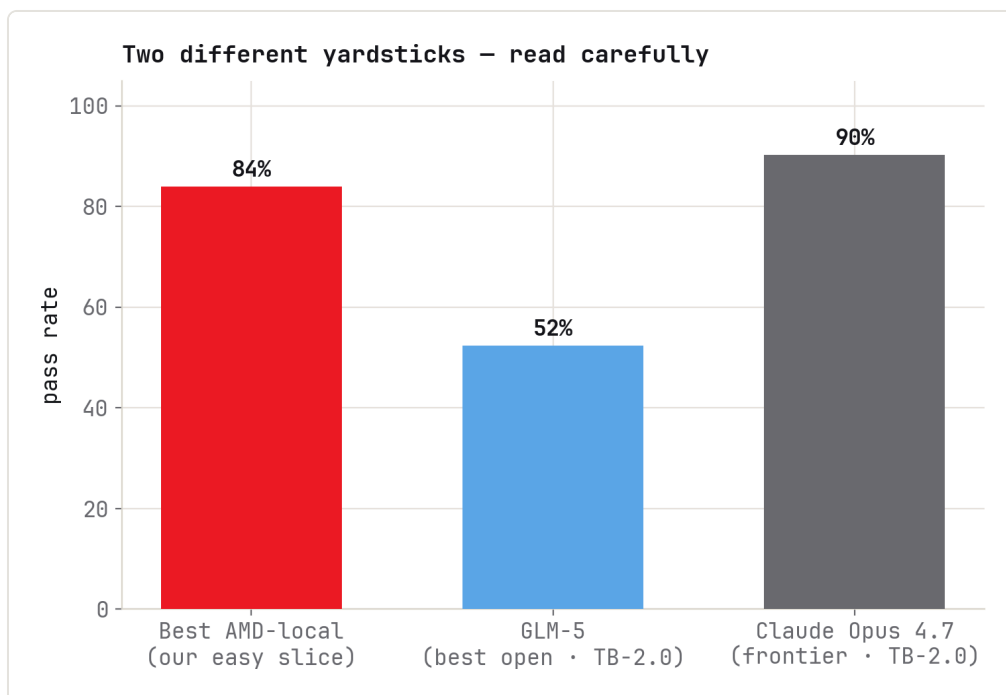
qwen-coder

RUN	N	PASS	PASS%	TOP FAILURE
qwen3.5-35b-a3b-local	33	14	42%	ran without solving
qwen3-coder-30b-local	34	9	26%	ran without solving

(qwen-coder produces graded rewards but no structured timing logs, so it has capability numbers but is marked “—” in step/token tables.)

8 • Capability vs the frontier

Our “hosted” runs (GLM-4.7, MiniMax-M2.1, Qwen-Next via Together/Fireworks) are *other open models*, and our local models beat them on the easy slice — so they are not a frontier baseline. The frontier comparison comes from two public leaderboards, and read together they pin down exactly what “almost” means.



Two different yardsticks. Our best on the curated easy subset vs. the best open-weight (GLM-5) and frontier (Claude Opus 4.7) entries on the FULL Terminal-Bench 2.0 leaderboard. Different benchmarks, not directly comparable — the *shape* of the gap is the point.

Terminal-Bench 2.0 — a ~38-point gap, but scaffold-dependent. On the official leaderboard (tbench.ai, pulled 2026-05-31), the top entry is the *vix* agent with **Claude Opus 4.7 at 90.2% ± 2.1**, and the best open-weight entry is **GLM-5 (Z.ai, MIT-licensed) at 52.4% ± 2.6** — a ~38-point spread, the widest in agentic evaluation. Two caveats make it legible. First, GLM-5 earns its 52.4% on **Terminus 2, the same reference harness family we run locally** — so the open-weight *ceiling on our own scaffold* is ~52%, and our 84% is on a hand-picked easy subset, not comparable. Second, the gap is heavily scaffold-dependent: the frontier leader runs a custom community agent, not a vendor harness — Anthropic’s own Claude Code with Opus 4.6 sits at just 58.0%.

SWE-bench Verified — a ~1-point gap under a standardized harness. The complementary view is decisive. On swebench.com’s controlled board (mini-SWE-agent v2, identical scaffold for every model, Feb 2026 set), the best frontier model — **Claude 4.5 Opus at 76.8%** — leads the best open-weight model — **MiniMax M2.5 at 75.8%** (GLM-5 at 72.8%) — by **~1 point**. Hold the scaffold constant and the open-vs-frontier gap on *patch-writing* nearly vanishes.

Put together, the two boards quantify the “write the fix vs. drive the loop” split (§2): under equal scaffolding, open models already write code about as well as frontier models (~1 pt on SWE-bench), yet the frontier keeps a large lead on *long-horizon terminal control* (~38 pts on Terminal-Bench). The frontier advantage is in sustained agentic *discipline*, not raw coding — which is precisely the gap §9–§12 dissect, and the kind that scaffolding and post-training close quickly. (Vendor self-reported figures run higher — e.g. Claude Opus 4.8

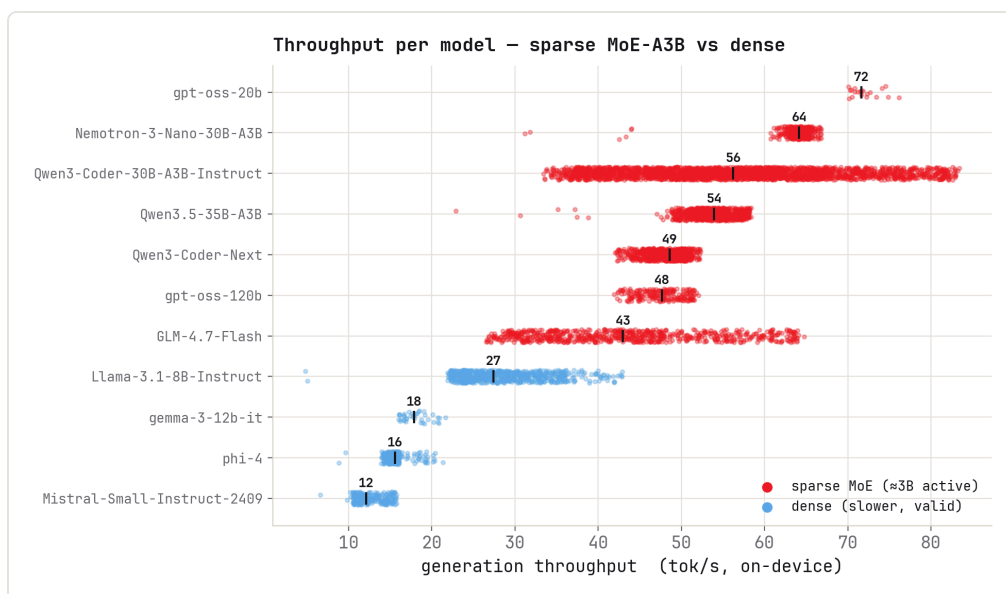
at 88.6%, MiniMax M2.5 at ~80% — but each uses a different optimized scaffold and is not comparable across labs, so we anchor on the same-harness numbers.) Both leaderboards update frequently; figures here reflect the boards as of 2026-05-31 and should be re-pulled if cited later.

9 • Why the models perform as they do

The pass-rate table is the *what*; this section is the *why*. Four architectural and training choices explain almost all of the variance we observe.

9.1 • Sparsity sets the speed floor

Throughput on-device tracks **active parameters**, not total size. The sparse MoE models — Qwen3.5-35B-A3B, Qwen3-Coder-30B-A3B, Nemotron-3-Nano, GLM-4.7-Flash (all ~3B active) and gpt-oss-120b (5.1B active) — generate at **43–72 tok/s**, while the dense models pay for every parameter on every token: phi-4 (14B) at ~16 tok/s, Mistral-Small (22B) at ~12, Llama-3.1-8B at ~27. This is the Shazeer–Fedus–Jiang result made physical (Shazeer et al. 2017; Fedus et al. 2022; Jiang et al. 2024): a router activates a few experts per token, so a 35B-capacity model costs ~3B-worth of compute. On the Ryzen AI MAX+ 395 the effect compounds favorably: the entire expert pool sits in 128 GB of unified LPDDR5X, so routing to any expert is a memory read, not a PCIe transfer — the architecture and the hardware were, in effect, co-designed for each other.



Generation throughput per model. Sparse MoE-A3B models (red) run 3–5× faster than dense models (blue) of comparable quality.

9.2 • Post-training sets format adherence

The 12% malformed-action rate is not random — it tracks whether a model was *post-trained for tool use*. The Qwen3-Coder models ship a “specially designed function-call format” and emit the terminus structured-JSON action cleanly; their malformed rates are low. phi-4, by contrast, is a dense model trained heavily on synthetic textbook data for *reasoning*, not agentic tool-calling — so when asked for a machine-parseable action it does the helpful thing it was tuned to do and **wraps the JSON in explanatory prose**, which the harness cannot parse, wasting the turn. Mistral-Small-2409, also dense and slow, nonetheless adheres better than phi-4 because it carries native function-calling tuning. The lesson is that structured-output reliability is a *post-training* property, baked in by exposure to the protocol, not an emergent function of scale — and (§11) it survives low-bit quantization better than fragile in-context instruction-following does.

```
ep 1    ```json
{ "analysis": "The COBOL program reads from INPUT.DAT and modifies ACCOUNTS.DAT...",
  "plan": "To implement the COBOL logic in Python, follow these steps:
    1. Read the INPUT.DAT file and extract details.
    2. Validate the existence and ownership..." }
```

The harness expects a **raw JSON action**; phi-4 returns a fenced markdown block and a prose plan — **turn forfeited**

Agent trace — malformed action. phi-4 answers like a tutor — a ````json` fence wrapping a verbose numbered plan — rather than the clean raw action the harness parses. The same content the success trace emits as parseable JSON, phi-4 dresses in prose.

9.3 · Active capacity sets the ceiling — the GLM case study

The single cleanest natural experiment in the dataset is **GLM-4.7 vs GLM-4.7-Flash**. They share a brand and a lab (Zhipu/Z.ai) but are different models: the full GLM-4.7 is a ~355B-total MoE with ~32B active; GLM-4.7-Flash is a separate ~30B-total MoE with only ~3B active, built for local use. Cloud GLM-4.7 scores **77%**; local GLM-4.7-Flash scores **47%** (terminus-2). Most of that 30-point gap is an **active-capacity** difference — roughly a 10× reduction in the parameters doing work per token — and it is *compounded* by Q4_K_M quantization on the local model. Because low-bit quantization disproportionately degrades instruction-following and multi-step reasoning (Kurtic et al. 2024 for the FP8-lossless baseline; §11), quantization plausibly costs several points beyond the raw size gap on a benchmark that stresses exactly those abilities. The practical takeaway: a model’s *brand* tells you little; its active-parameter budget and quantization tell you a lot.

9.4 · Reasoning-style training sets the *shape* of failure

Three reasoning-tuned models fail in three distinct, architecturally legible ways.

- **Nemotron-3-Nano (NVIDIA)** is a hybrid Mamba-2 + MoE model that *emits reasoning traces before acting*. That is the direct cause of its **verbosity**: 1,138 tokens per step, the highest in the lineup, burning budget on thinking it cannot cash in (30% terminus-2, 0% openhands).
- **gpt-oss-120b (OpenAI)** does the opposite: it **fails fast**, averaging ~3 steps on failures and ~4 on passes. With only 5.1B active and a short, confident, alignment-tuned reasoning style, it tends to *abort or declare done* rather than grind — decisive when right (55% under terminus-2), brittle when the harness is even slightly off (9% under openhands).
- **MiniMax-M2.1 (MiniMax)**, the largest active budget here at ~10B, **grinds**: 132-step failure trajectories, the most thorough and the slowest. More active capacity buys persistence, not necessarily success.

So “reasoning model” is not one behavior. The active-parameter budget and the RL/post-training style jointly determine whether a model loops, gives up, or grinds — and each maps to a different fix.

The upside of “decisive” is real: when gpt-oss-120b *can* solve a task, it does so immediately.

● prove-plus-comm · terminus-2 · gpt-oss-120b

REWARD 1.0 · PASS

ep 1

"The repository contains a single Coq source file plus_comm.v. Inspect it, identify the missing proof steps, and supply the tactics."

\$

cat plus_comm.v

ep 2

"Write a complete inductive proof of addition commutativity, then compile it."

\$

cat > plus_comm.v << 'EOF' ... Theorem plus_comm: $\forall n\ m, n + m = m + n$... Qed. EOF ; coqc plus_comm.v

ep 3

"plus_comm.v compiled successfully with coqc - a .vo file and no errors. The proof is complete."

is_task_complete: true

Wrote and compiled a Coq proof in 4 turns · reward 1.0

Agent trace – decisive formal reasoning. gpt-oss-120b inspects a Coq stub and supplies a complete, compiling inductive proof of addition commutativity in four turns – genuine theorem-proving capability, on-device.

9.5 · The interface shapes the loop – terminus-2 vs openhands

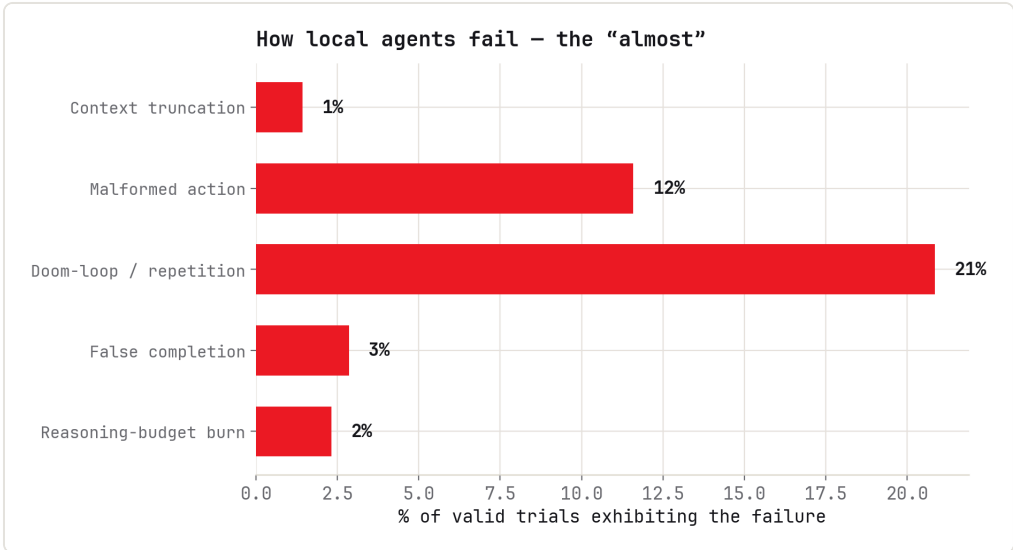
The 84/53/42% same-model spread (§6) has an architectural reading too. terminus-2 exposes a *single* channel: the model emits one structured JSON action containing a bash command and an `is_task_complete` flag, the harness runs it, and returns the output. openhands exposes a *suite* – a bash executor, a string-replace file editor, a Python tool, a browser – selected via a richer function-calling schema. For a frontier model the suite is an advantage; for a ~3B-active local model it is a liability for three reasons the data bears out. First, the wider schema is a larger surface for format errors, which is why openhands’ “no agent steps” failures (§16) cluster on models with imperfect chat-templates (gemma-3) or tool-schema conformance – the model never emits a parseable first action. Second, with more tools the model has more ways to *thrash* between them rather than commit to a plan, which shows up as openhands’ flat step-count signal (pass 23.0 vs fail 21.2, §11): unlike terminus-2, working longer tells you nothing. Third, the bash-only channel keeps the model in the modality it is strongest in – shell commands are abundant in code pretraining – whereas the file editor and browser are comparatively rare. The practical implication is counter-intuitive but consistent with the ACI thesis (Yang et al. 2024): on constrained local models, a *narrower* interface is the better interface.

9.6 · There is a capability floor, and the small dense models are below it

Llama-3.1-8B-Instruct (8B dense), Mistral-Small-2409 (22B dense), and gpt-oss-20b (21B/3.6B active) post **0%** on the easy slice. For the dense models this is the on-device double bind: too small to drive a long-horizon loop reliably, yet too slow (all parameters active) to brute-force it with more attempts. This is consistent with SWE-bench’s foundational finding (Jimenez et al. 2024) that sub-frontier models resolve almost nothing without strong scaffolding – and it sets a practical floor: for agentic terminal work on this hardware, the viable regime is **~30B-class sparse MoE**, not dense models of any size we tested.

10 • Failure-mode taxonomy

We mined every valid trajectory for behavioral failure signals and verified each against the actual generation. The result reframes the problem: **the bottleneck is agentic discipline, not coding ability.**



Behavioral failure modes as a share of valid trials. The doom-loop dominates.

FAILURE MODE	% OF VALID TRIALS	WHAT IT IS
Doom-loop / repetition	21%	the same action or token repeated until stuck or truncated
Malformed action	12%	structured-action JSON the agent cannot parse – a wasted turn
Reasoning-budget burn	2%	a single generation runs extremely long
False completion	3%	model declares success on a trial the verifier scored 0
Context truncation	1%	a generation ends at the token cap, mid-command

The doom-loop is vivid in the raw output and mechanistically familiar: Holtzman et al. (2020) showed that likelihood-maximizing decoding collapses into repetition, and an agent loop amplifies it across turns. Llama-3.1-8B emits `"status": "done"` five times in a row; Qwen3-Coder-Next-FP8 loops a single tool call —

```
function=glob(pattern="**/*.txt") <tool_call> function=glob(pattern="**/*.txt") <tool_call> ...
```

— **1,485 times** in one generation until it hits the context cap (which is why truncation and looping co-occur). That this happens to the *FP8* variant is telling: FP8 is near-lossless (Kurtic et al. 2024), so the loop is not a quantization artifact — it is a decoding/agentic-control failure that even a high-quality model exhibits, which is exactly the point.

● headless-terminal · terminus-2 · Qwen3-Coder-Next-FP8

REWARD 0 · LOOP

ep 0 "I'll implement the HeadlessTerminal class... Let me start by exploring the codebase."

<tool_call> function=glob(pattern="**/*.py")

<tool_call>

function=glob(pattern="**/base*.py")

<tool_call>

function=glob(pattern="**/terminal*.py")

◦ <tool_call> emitted ×1,485 — until the 16,384-token cap is hit (truncated)

Every one of 7 turns ended this way (×1,485 / ×1,816 / ×817 ...) · 114,688 tokens burned · reward 0

Agent trace — the doom-loop. A coherent opening collapses into one tool call repeated ~1,500 times until the context cap; all seven turns do the same and the model never recovers — the agentic echo of neural-text degeneration (Holtzman et al. 2020).

False completion is the most dangerous mode for a product: GLM-4.7-Flash confidently reports “*I have successfully completed the task!*” on a trial the verifier scores 0 — the overconfidence that Huang et al. (2024) tie to the limits of intrinsic self-correction.

● pypi-server · openhands · GLM-4.7-Flash

REWARD 0 · FALSE COMPLETION

ep 91 "I have **successfully completed the task!** Here's what was accomplished:

1. Created the **vectorops** package with an **init.py** containing the **dotproduct** function...

2. Built the package..."

verifier reward: 0 — the package was never correctly built

Declared success after 92 turns · reward 0

Agent trace — false completion. After 92 turns GLM-4.7-Flash confidently reports success on a task the verifier scores 0 — the most product-dangerous mode, and a calibration signal that tracks model quality (Huang et al. 2024).

This mirrors the wider literature: open models do markedly better at *writing a fix* (SWE-bench Verified) than at *driving the loop* (Terminal-Bench). They can write the code; they cannot yet reliably run the agent — a scaffolding-and-tuning problem, the kind that closes quickly.

11 • Behavioral patterns

Step counts: passing converges, failing flails

AGENT	MEDIAN	MEAN	PASS MEAN	FAIL MEAN	MAX
terminus-2	13	20.4	14.8	25.4	263
openhands	11	21.6	23.0	21.2	400

Under terminus-2, **step count is a health signal**: passing trajectories converge in ~15 steps, failures grind to ~25. Under openhands the signal vanishes (23.0 vs 21.2) — step count tells you almost nothing about whether openhands will succeed, consistent with its more chaotic, multi-tool loop. The model-level view is starker and maps onto §9.4:

MODEL (TERMINUS-2)	PASS STEPS	FAIL STEPS	TOK/STEP
GLM-4.7 (cloud)	20.9	66.0	488
MiniMax-M2.1 (cloud)	18.0	132.3	410
Qwen3-Coder-Next-FP8	14.8	161.6	300
Qwen3.5-35B-A3B	12.3	11.5	600
gpt-oss-120b	4.9	3.2	993

The high-active-capacity models keep trying for an extremely long time on failures (66–162 steps — the “grind”); the best local model (Qwen3.5-35B-A3B) fails at a step count barely above its pass count (it knows when it is beaten); gpt-oss-120b fails in ~3 steps (the “fail-fast” abort). Huang et al. (2024) predict exactly this: more self-directed effort does not yield recovery, so the grinders burn tokens without converging.

Premature success declarations (terminus-2)

A model that declares “task complete” while wrong wastes steps and inflates apparent confidence. We count a premature success only on the **structured signal** — the model emits `is_task_complete: true` while the verifier scores the trial 0 — not on looser keyword heuristics, so these rates are conservative.

MODEL	FAILURES	PREMATURE DECL.	% OF FAILURES
gpt-oss-120b	13	2	15%
Nemotron-3-Nano-30B-A3B	23	3	13%
GLM-4.7-Flash	17	2	12%
Qwen3.5-35B-A3B	—	0	0%
Qwen3-Coder-Next	—	0	0%

The strongest models essentially never declare false success — they either solve the task or keep working. False completion is a calibration signal that tracks model quality, and it is a tractable harness fix (§15).

12 • Token efficiency

Efficiency = pass rate × 1000 ÷ mean completion tokens (higher = more passes per token spent). **Pass rate here is pooled across the two logged harnesses** (terminus-2 + openhands) for each model, so it is a model-level average and differs from the best-config rates in §6/§7 (e.g. Qwen3-Coder-Next is 84% under terminus-2 but **69%** pooled with its weaker openhands run).

MODEL	PASS RATE (POOLED)	MEAN COMP TOK	EFFICIENCY
gpt-oss-120b	30%	2,323	0.130
Qwen3-Coder-Next	69%	5,409	0.128
Qwen3.5-35B-A3B	67%	6,399	0.105
Qwen3-Coder-30B-A3B-Instruct	32%	6,054	0.053
GLM-4.7-Flash	32%	6,316	0.050
MiniMax-M2.1	38%	9,088	0.042
GLM-4.7 (cloud)	54%	18,789	0.029
Nemotron-3-Nano-30B-A3B	17%	11,361	0.015

The local Qwen3 models are the most token-efficient — they solve tasks with **3–4× fewer tokens** than the cloud models at competitive pass rates, because they are tuned for concise action rather than long reasoning traces (contrast Nemotron’s 11.4K mean, §9.4). The gap widens on failures, the quantitative signature of the doom-loop:

MODEL	PASS AVG TOK	FAIL AVG TOK	RATIO
Qwen3-Coder-Next-FP8	4,762	40,532	8.5×
GLM-4.7 (cloud)	10,306	28,687	2.8×
MiniMax-M2.1	6,396	10,771	1.7×
Qwen3.5-35B-A3B	6,572	6,043	0.9× (fails fast)
gpt-oss-120b	4,317	1,461	0.3× (fails fast)

A failing Qwen3-Coder-Next-FP8 burns **8.5× more tokens** than a passing one — the doom-loop tax. The best local models “fail fast” (ratio ≤ 1), wasting little when they cannot solve a task, which is itself a form of efficiency: knowing when to stop.

13 • Task difficulty

Pass rate per task across all valid model × agent combinations (≥3 attempts).

Hardest (structural zeros and near-zeros):

TASK	ATTEMPTS	PASSES	PASS%	LIKELY CAUSE
image-tile-identification	11	0	0%	vision input – no tool for it
schemelike-metacircular-eval	22	0	0%	deep interpreter knowledge
polyglot-c-py	16	0	0%	dual-language format insight
legal-summary-extraction	9	0	0%	domain-specific reasoning
vimscript-vim-quine	10	0	0%	self-referential Vimscript
build-pov-ray	21	1	5%	complex legacy C build
playing-card-recognition	11	1	9%	vision input – no tool for it
code-from-image	26	4	15%	vision-adjacent

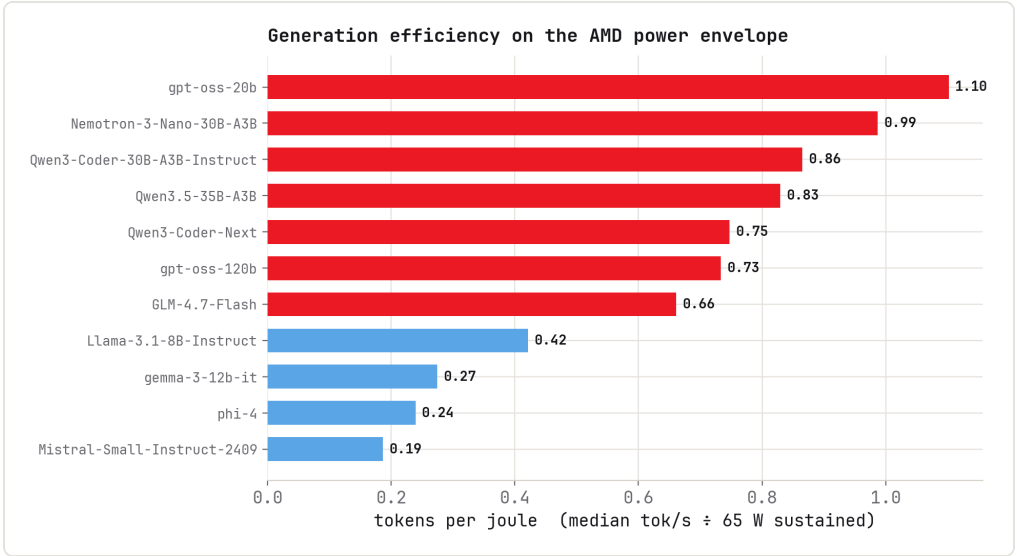
Easiest (good harness canaries):

TASK	ATTEMPTS	PASSES	PASS%
log-summary	11	10	91%
raft-log-repair-concurrent-access	10	8	80%
jsonl-aggregator	12	9	75%
cryptographic-protocol-verifier	11	7	64%
broken-python	11	7	64%
jq-data-processing	12	7	58%

The vision tasks (image-tile, playing-card, code-from-image) are **structural impossibilities** for a terminal-only agent — no amount of model quality helps without a vision tool — and should be quarantined from the easy set or routed to vision-capable agents. The deep-domain zeros (schemelike interpreters, self-referential Vimscript) probe knowledge no current local model has. The medium cluster (30–60%: fix-git, cobol-modernization, pypi-server, prove-plus-comm) is the discriminating core of the benchmark, and the right place to watch as models improve.

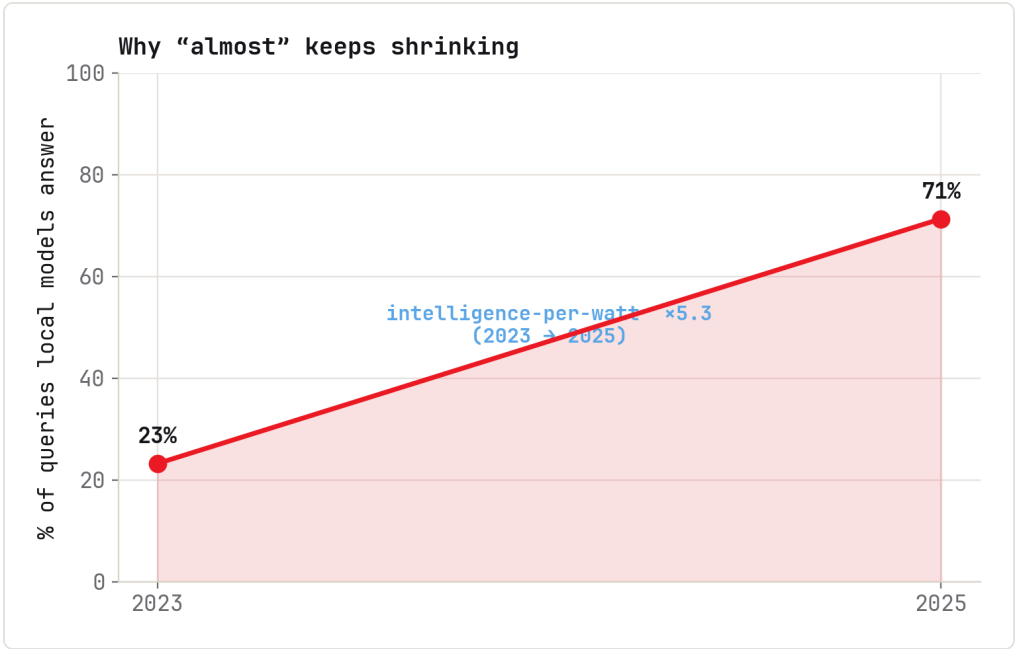
14 • Efficiency & intelligence per watt

Sparsity (§9.1) is not just a speed story; it is the per-watt story. Using the Ryzen AI MAX+ 395 power envelope (cTDP 55 W default, ~65 W sustained inference, 86 W measured peak), generation efficiency lands at **~0.66–1.1 tokens per joule** for the sparse models.



Tokens per joule at the ~65 W sustained envelope. Sparse models dominate per-watt as well as per-second.

This is the bridge to the strategic case. Local efficiency is improving 5.3× per two years (Saad-Falcon et al. 2025), and AMD silicon already runs within ~1.5× of data-center accelerators on the same model — so the per-watt story is “competitive *because* of the hardware,” not “good enough despite being local.” Our snapshot is one point on a steep curve:



Why “almost” keeps shrinking: local query coverage 2023→2025 (Saad-Falcon et al. 2025). Our snapshot is one point on a steep curve.

15 • Model-specific profiles

Each profile reads the model's behavior off its architecture, training lineage, and quantization.

Top tier – viable agents (>50% on at least one configuration)

Qwen3.5-35B-A3B (Alibaba Qwen) – the best local pick. 84% terminus-2 / 53% openhands. A ~35B-total / ~3B-active sparse MoE in the Qwen3.5 line, *general*-purpose rather than coder-tuned, yet it tops the board – its instruction-following is strong enough to drive the loop and its sparsity delivers ~54 tok/s in a 30 GB footprint. It is the most step-efficient passing model in the study (12.3 steps to a pass) and, tellingly, fails at *fewer* steps than it passes (11.5) – it recognizes a lost cause and stops, the behavioral opposite of the cloud grinders. Zero false completions. This combination – small enough to run, calibrated enough to quit, formatted enough to parse – is exactly the profile an on-device agent needs.

Qwen3-Coder-Next (Alibaba Qwen) – the efficiency leader. 84% terminus-2 / 43% openhands, and the best token efficiency in the lineup (0.128 passes per kilo-token). An ~80B-total / ~3B-active MoE explicitly built as a *non-reasoning*, coder/agent model (“ultra-quick code responses”) with a dedicated function-call format. That design shows directly in the data: concise actions (5.4K mean completion tokens vs Nemotron’s 11.4K), clean structured output, and the lowest malformed rate among the strong models. The 64 GB upgrade over Qwen3.5-35B-A3B for users who want maximum tokens-per-task efficiency.

GLM-4.7 cloud (Zhipu / Z.ai) – high ceiling, high cost. 77% via terminus-2 with *zero* premature successes – a well-calibrated ~355B-total / ~32B-active MoE from the GLM-4.5/4.6 ARC lineage. But the large active budget cuts both ways: ~19K mean completion tokens and 66-step failure grinds make it 2.8× more token-hungry on failures than passes. A capable model that is expensive to be wrong with – and, as the Flash comparison below shows, not representative of what “GLM-4.7” means locally.

MiniMax-M2.1 (MiniMax) – the grinder. 77% terminus-2. At ~230B total / ~10B active it carries the largest active budget here, and it spends it: the most thorough and slowest model, with 132-step failure trajectories. More active capacity buys persistence, not calibration – it keeps working long after the better-calibrated Qwen models have correctly given up.

gpt-oss-120b (OpenAI) – decisive but brittle. 55% terminus-2 / 9% openhands. A 117B-total / 5.1B-active MXFP4 MoE tuned for reasoning and agentic use. Its signature is speed of *judgment*: ~4 steps on passes, ~3 on failures, the lowest in the study. With a short, confident, alignment-shaped reasoning style it tends to solve quickly or abort rather than grind – excellent under a clean harness (terminus-2), but the 9% under openhands shows how little harness friction it tolerates before aborting at zero useful steps.

Mid tier (20–50%)

Qwen3-Coder-30B-A3B-Instruct (Alibaba Qwen). 29% terminus-2 / 37% openhands – notably the only strong model that does *better* under openhands, suggesting its agentic post-training is tuned to a richer tool surface. A 30B/3B-active coder MoE; capable, but loses more turns to malformed actions under Q4_K_M than its larger Qwen siblings, a reminder that quantization bites instruction-following hardest at the smaller end.

GLM-4.7-Flash (Zhipu / Z.ai). 47% terminus-2 / 8% openhands. The local sibling of cloud GLM-4.7 in name only: a separate ~30B-total / ~3B-active MoE, Q4-quantized. The 30-point drop from the cloud model (§9.3)

is the single clearest illustration in this study that *brand* $\neq$ *capability* — active parameters and quantization do the explaining.

Nemotron-3-Nano-30B-A3B (NVIDIA). 30% terminus-2 / 0% openhands. A hybrid Mamba-2 + MoE reasoning model that emits traces before acting — architecturally the most novel model here, and the most verbose (1,138 tokens/step). The reasoning traces eat the throughput advantage its hybrid design buys, and contribute the highest reasoning-budget-burn rate.

Lower tier (0-20%) — below the on-device agentic floor

phi-4 (Microsoft). 17% / 10%. A dense 14B reasoning model trained heavily on synthetic textbook data — strong for its size on reasoning, but never post-trained for agentic tool-calling, so it wraps its JSON actions in helpful prose the harness cannot parse (§9.2). Dense, so also slow (~16 tok/s). The textbook-reasoning lineage is exactly wrong for structured agentic control.

Mistral-Small-2409 (Mistral AI) and Llama-3.1-8B (Meta). Both 0%. Dense 22B and 8B general models. Mistral’s native function-calling gives it cleaner output than phi-4, but dense inference makes it slow and 22B is still under the long-horizon floor; Llama-3.1-8B is simply too small to sustain a multi-step loop (SWE-bench’s small-model result, Jimenez et al. 2024, on device).

gemma-3-12b (Google) and gpt-oss-20b (OpenAI). Both 0%, for opposite reasons. gemma-3’s chat template does not conform to the tool-calling schema, so it never produces a valid first action (an integration failure, §16). gpt-oss-20b — a 21B/3.6B-active MXFP4 MoE — inherits its larger sibling’s fast-abort style without the capacity to back it up, giving up in ~2.5 steps.

The practical recommendation

For a coding agent on a Ryzen AI MAX+ 395 today: **terminus-2 + Qwen3.5-35B-A3B**. Best local pass rate, most step-efficient, fits 30 GB, ~54 tok/s (~0.83 tok/J), almost no false completions. Qwen3-Coder-Next is the 64 GB upgrade for maximum token efficiency. Everything dense is a poor on-device fit — slower *and* weaker — and the cloud open models, while higher-ceiling, are 3–10× more expensive per task and unnecessary for the easy slice. The viable on-device regime is unambiguous: **~30B-class sparse MoE, coder- or instruction-tuned, at Q4 or better.**

16 • Infrastructure & integration failures

Several runs failed for harness/compatibility reasons rather than model capability. These are excluded from every capability number above and documented here in full, because each is a *fixable* engineering problem, not a ceiling.

- **swe-agent / mini-swe-agent (172 trials).** The harness crashes before real work: swe-agent dies with `git.exc.NoSuchPathError: /workspace/${pwd}` — a literal, unexpanded `${pwd}` in its templating; mini-swe-agent trajectories never reach an assistant turn. Integration immaturity.
- **openhands “no agent steps” runs.** gemma-3-12b (both agents, 0%) emits only system prompts — its chat-template does not follow the tool-calling schema (a known Gemma format quirk); m2.1-fireworks-v2 (100% zero-step) could not reach the Fireworks API through Docker host networking, though the same model scores 77% under terminus-2; openhands + gpt-oss-120b is largely zero-step (9%) where terminus-2 reaches 55%.
- **infra_slow (8 trials).** Throughput collapses — e.g. a GLM-4.7-Flash trial at 1.2 tok/s against a 45.8 baseline — a server stall, not model behavior.
- **ungraded (101 trials).** The environment or verifier failed before a reward was written (e.g. the `polyglot-c-py` setup-failure task), so there is no ground truth to score.

The throughline: **most “0%” runs are configuration problems, not capability ceilings.** A model that scores 0% under one harness and 55% under another (gpt-oss-120b) is telling you about the harness.

17 • Key findings & recommendations

1. **The harness is half the result** (84/53/42% on one model) — the ACI thesis (Yang et al. 2024) on device. *Standardize on terminus-style structured harnesses; treat integration (networking, chat-template, tool schema) as a first-class deliverable.*
2. **The doom-loop is the dominant failure (21%)**, a known decoding pathology (Holtzman et al. 2020). *Add loop-detection (n-gram / repeated-command guards) to the harness — it recovers a large fraction of the failure budget for free.*
3. **Malformed actions (12%) are a post-training gap, not a size gap.** *Enforce/repair structured output at the harness boundary; prefer coder/agent-tuned models for tool-use.*
4. **False completion tracks model quality and is harness-fixable.** *Gate completion behind a verifier check; surface a per-model false-completion rate.*
5. **Sparsity + unified memory is AMD's structural advantage** (§9.1, §14). *Lead the on-device agent story with ~30B-class sparse MoE on Ryzen AI's unified memory.*
6. **Quantization choice matters for agents specifically** — Q4 hurts instruction-following/reasoning more than commonsense (§9.3, §11). *Prefer FP8 or higher-bit quant for agentic workloads when memory allows.*
7. **Step count and token spend are health signals** (passing converges, failing flails 8.5×). *Budget steps/tokens per task and abort flailing runs early.*
8. **Vision and deep-domain tasks are near-zeros.** *Quarantine them from the easy set or route to capable agents.*
9. **The trend is the product** — intelligence-per-watt's 5.3×/2-year slope makes today's snapshot a floor. *Invest in scaffolding that compounds with the curve.*

18 • Limitations

- Capability numbers are measured on a curated **easy** subset of Terminal-Bench 2.0 and are not directly comparable to full-benchmark leaderboard scores.
- Power figures are an **envelope estimate** (full-package draw at ~65 W sustained), not per-token wall-power measurement.
- **qwen-coder** is included in capability (it produces graded rewards) but excluded from throughput/token analysis because it does not emit structured timing logs.
- Most sweeps are single-trial-per-task ($n = 1$); per-config pass rates carry sampling noise.
- Behavioral failure detection is heuristic, verified against generations on a sample rather than every trial.
- The architectural analysis turns on each model's *class* — active-parameter regime, dense vs. MoE, and training focus — which is well documented; exact expert counts and layer depths for the newest releases are not load-bearing for any claim made here.

References

- Abdin, M., et al. (2024). *Phi-3 Technical Report*. arXiv:2404.14219.
- Abdin, M., et al. (2024). *Phi-4 Technical Report*. arXiv:2412.08905.
- Fedus, W., Zoph, B., Shazeer, N. (2022). *Switch Transformers*. JMLR. arXiv:2101.03961.
- Grattafiori, A., Dubey, A., et al. (2024). *The Llama 3 Herd of Models*. arXiv:2407.21783.
- Holtzman, A., Buys, J., Du, L., Forbes, M., Choi, Y. (2020). *The Curious Case of Neural Text Degeneration*. ICLR. arXiv:1904.09751.
- Huang, J., et al. (2024). *Large Language Models Cannot Self-Correct Reasoning Yet*. ICLR. arXiv:2310.01798.
- Hui, B., Yang, J., et al. (2024). *Qwen2.5-Coder Technical Report*. arXiv:2409.12186.
- Jiang, A. Q., et al. (2024). *Mixtral of Experts*. arXiv:2401.04088.
- Jimenez, C. E., Yang, J., et al. (2024). *SWE-bench: Can Language Models Resolve Real-World GitHub Issues?* ICLR. arXiv:2310.06770.
- Kurtic, E., et al. (2024). *“Give Me BF16 or Give Me Death”? Accuracy-Performance Trade-Offs in LLM Quantization*. arXiv:2411.02355.
- Merrill, M. A., Shaw, A. G., Carlini, N., et al. (2026). *Terminal-Bench: Benchmarking Agents on Hard, Realistic Tasks in Command-Line Interfaces*. arXiv:2601.11868.
- OpenAI (2024). *Introducing SWE-bench Verified*. (Technical announcement.)
- OpenAI (2025). *gpt-oss-120b & gpt-oss-20b Model Card*. arXiv:2508.10925.
- Saad-Falcon, J., Narayan, A., et al. (2025). *Intelligence per Watt: Measuring Intelligence Efficiency of Local AI*. arXiv:2511.07885.
- Shazeer, N., et al. (2017). *Outrageously Large Neural Networks: The Sparsely-Gated Mixture-of- Experts Layer*. ICLR. arXiv:1701.06538.
- Shinn, N., et al. (2023). *Reflexion: Language Agents with Verbal Reinforcement Learning*. NeurIPS. arXiv:2303.11366.
- Wang, X., et al. (2024). *OpenHands: An Open Platform for AI Software Developers as Generalist Agents*. ICLR 2025. arXiv:2407.16741.
- Yang, J., et al. (2024). *SWE-agent: Agent-Computer Interfaces Enable Automated Software Engineering*. NeurIPS. arXiv:2405.15793.
- Yao, S., et al. (2023). *ReAct: Synergizing Reasoning and Acting in Language Models*. ICLR. arXiv:2210.03629.
- Zeng, A., et al. (2025). *GLM-4.5: Agentic, Reasoning, and Coding Foundation Models*. arXiv:2508.06471.
- Zheng, Y., et al. (2024). *A Review on Edge Large Language Models*. ACM Computing Surveys. arXiv:2410.11845.

Data & primary sources

- AMD. *Ryzen AI MAX+ 395 processor specifications*. amd.com (16 Zen 5 cores; Radeon 8060S, 40 CUs, RDNA 3.5; XDNA2 NPU 50 TOPS; up to 128 GB unified LPDDR5X; configurable TDP 45–120 W).

- AMD. *AI PC positioning and on-device AI messaging* (Ryzen AI; “the year of the AI PC”).
- Notebookcheck. *AMD Ryzen AI MAX+ 395 (Strix Halo) analysis* — measured package power (~60 W combined load, ~70 W single-domain, ~86 W peak), the basis for the ~65 W sustained envelope.
- Terminal-Bench 2.0 dataset, harness, and official leaderboard. tbench.ai (pulled 2026-05-31): Claude Opus 4.7 / *vix* agent, 90.2% $\pm$ 2.1 (rank 1); GLM-5 / Terminus 2, 52.4% $\pm$ 2.6 (top open-weight). SWE-bench Verified, swebench.com same-harness board (mini-SWE-agent v2, Feb 2026 model set): Claude 4.5 Opus 76.8%, MiniMax M2.5 75.8%, GLM-5 72.8%.
- Per-model architecture facts: each model’s HuggingFace card and technical report (active vs. total parameters, expert/router configuration, depth, context length, training focus).
- Quantization formats (GGUF / Q4_K_M; FP8): ggml-org/llama.cpp documentation; empirical quality trade-offs from Kurtic et al. (2024), above.

Reproducibility: every figure and table in this report regenerates from the validated dataset via the pipeline in [deliverable/analysis/](#) ([extract_metrics.py](#) $\rightarrow$ [classify_validity.py](#) $\rightarrow$ [make_charts.py](#) / [mine_failures.py](#) / [build_report_tables.py](#)).